

Unsupervised Learning

Find groups in the data

- No labels nor response -> unsupervised
- Define groups based on similarity



Group customers, target ads

- A priori, you can't really put labels on customers
- Group similar customers

You can then try to interpret the grouping, and send targeted ads to the groups


Note: you may want to assign labels to groups a posteriori

Defining similarity

After a pre-processing step, you have a data matrix with **n** rows (observations) and **p** columns (features). Each row is a “point”.

- How to define similarity between points?
- If the features are numerical, we can use the Euclidean distance
- What if some features are categorical?
 - Ignore
 - Embed into numerical

Euclidean distance




$$d(a, b)^2 = \sum_{i=1:2} (x_i^a - x_i^b)^2$$


Can be generalised from 2-D to n-D

KMeans


K-means




K-means




K-means




K-means




K-means




K-means




K-means



K-means



K-means



K-means - Summary

- Start with K “means” drawn at random
- Assign data points to the nearest one
- Update the position of the means to correspond to the mean of those points
- Repeat...

K-means - Pros and cons

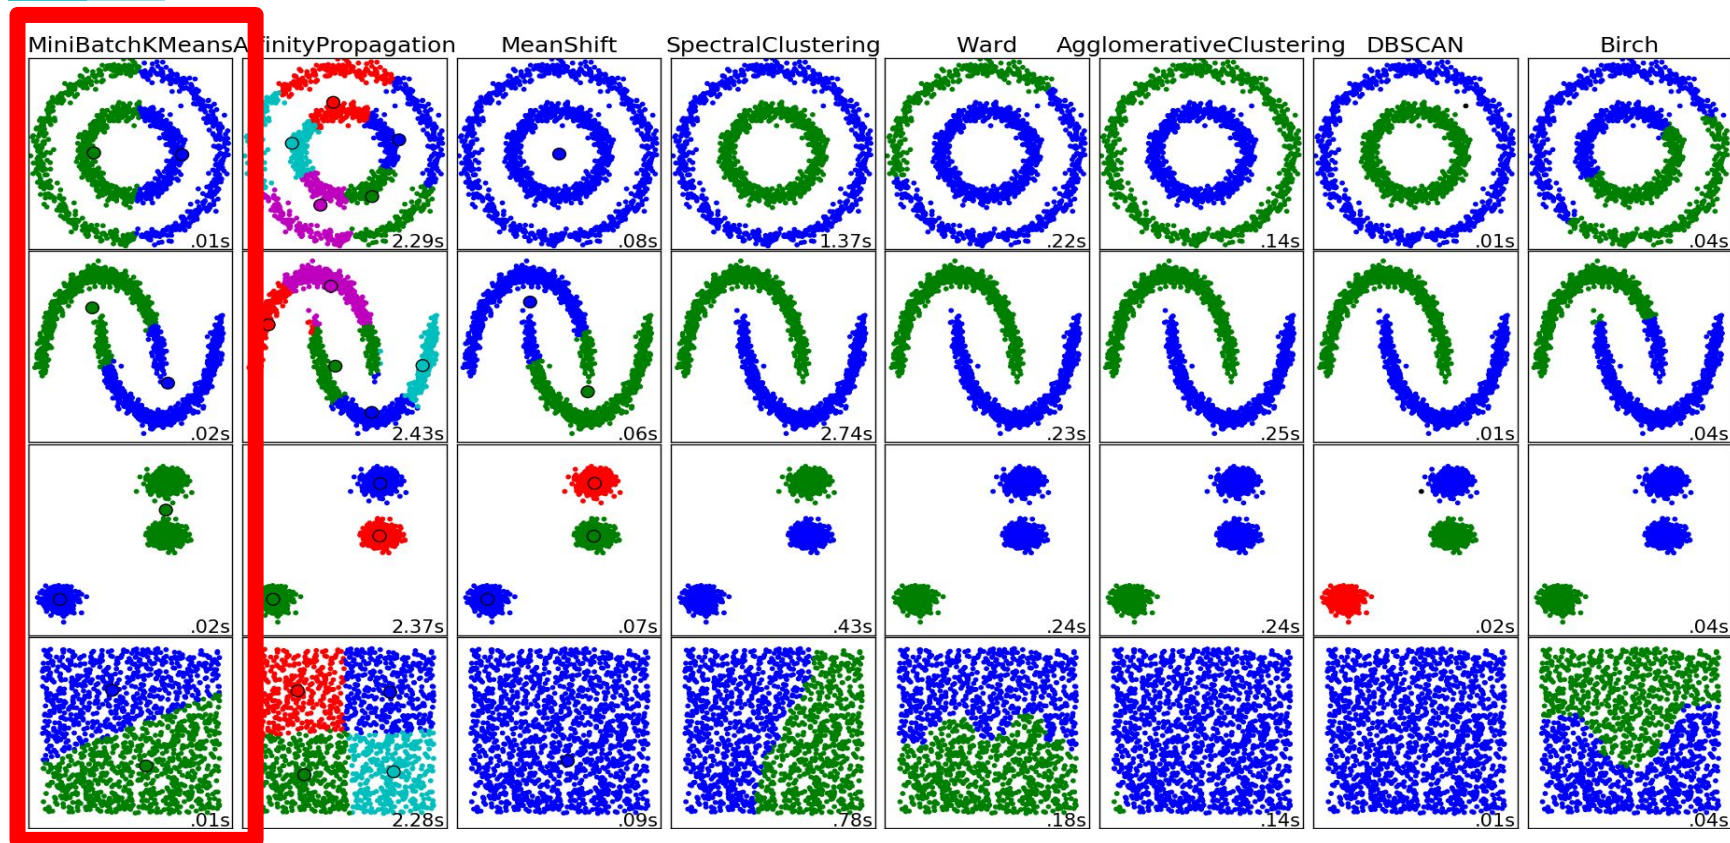
Pros

- Cheap to compute
- Easy to interpret
- Efficient implementations available
- Assigning a new point is straightforward

Cons

- Need to guess K
- Clusters are globular
- Sensitive to initialisation
- Sensitive to noise

Comparisons of clustering algorithms





Hands-on session

01-unsupervised_learning.ipynb

Hierarchical Clustering


Hierarchical clustering

Building a hierarchy of clusters sequentially.


- **Agglomerative** (bottom-up): Start considering each point as a cluster then merge the closest ones and repeat
- **Divisive** (top-down): Start with one single cluster and divide to have groups with reduced variance

Let's take a look at agglomerative hierarchical clustering (a.k.a. linkage)


Linkage algorithm




Linkage algorithm




Linkage algorithm




Linkage algorithm



Linkage algorithm





Linkage algorithm



Linkage algorithm

Two strategies to merge clusters:

- **Single linkage:** closest point distance (build spanning trees)
- **Complete linkage:** furthest point distance (to avoid elongated clusters)



Linkage: pros and cons

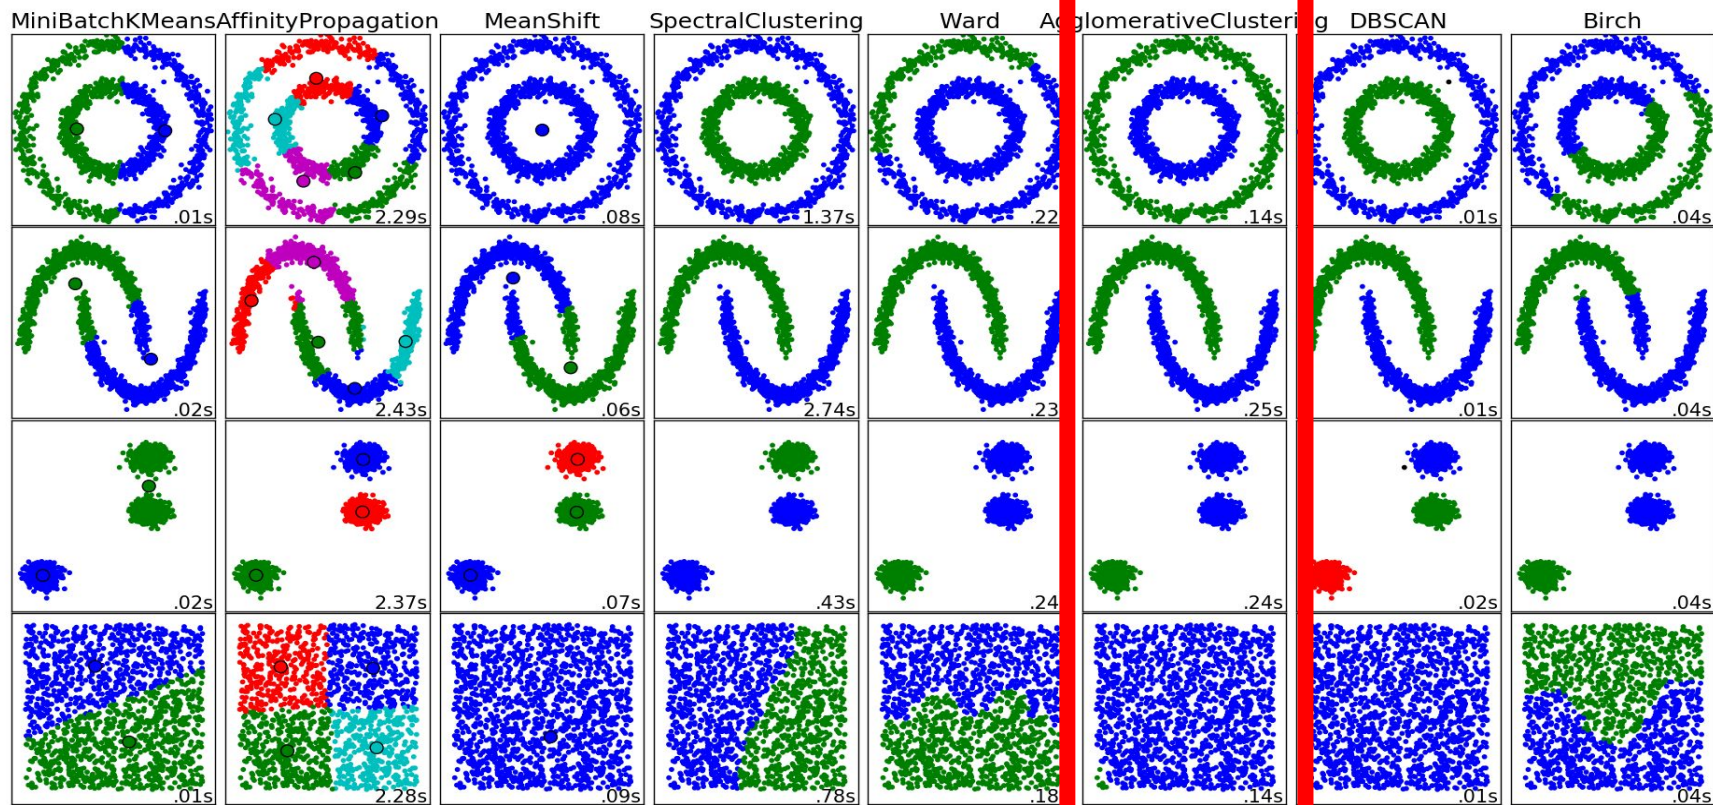
Pros

- Clusters are not necessarily globular
- No dependence upon initialisation
- Dendrogram shows a good summary

Cons

- Slower than K-means
- Still need to pick a number of clusters
- Assigning a new point is not straightforward
- Sensitive to noise

Comparisons of clustering algorithms





Hands-on session

01-unsupervised_learning.ipynb

DBSCAN

DBSCAN

Clusters = Zones of **high-density**


Two parameters: **min_samples** and **eps**

Algorithm:


- Start at a random point, consider all points within radius **eps**
- If that covers **min_samples**, keep that ball
 - Expand by considering esp-balls around every point of the current ball and iterate
- Otherwise mark the point as **noise**

... let's see this in action


DBSCAN




DBSCAN




DBSCAN




DBSCAN




DBSCAN




DBSCAN




DBSCAN




DBSCAN



DBSCAN



DBSCAN



DBSCAN: pros and cons

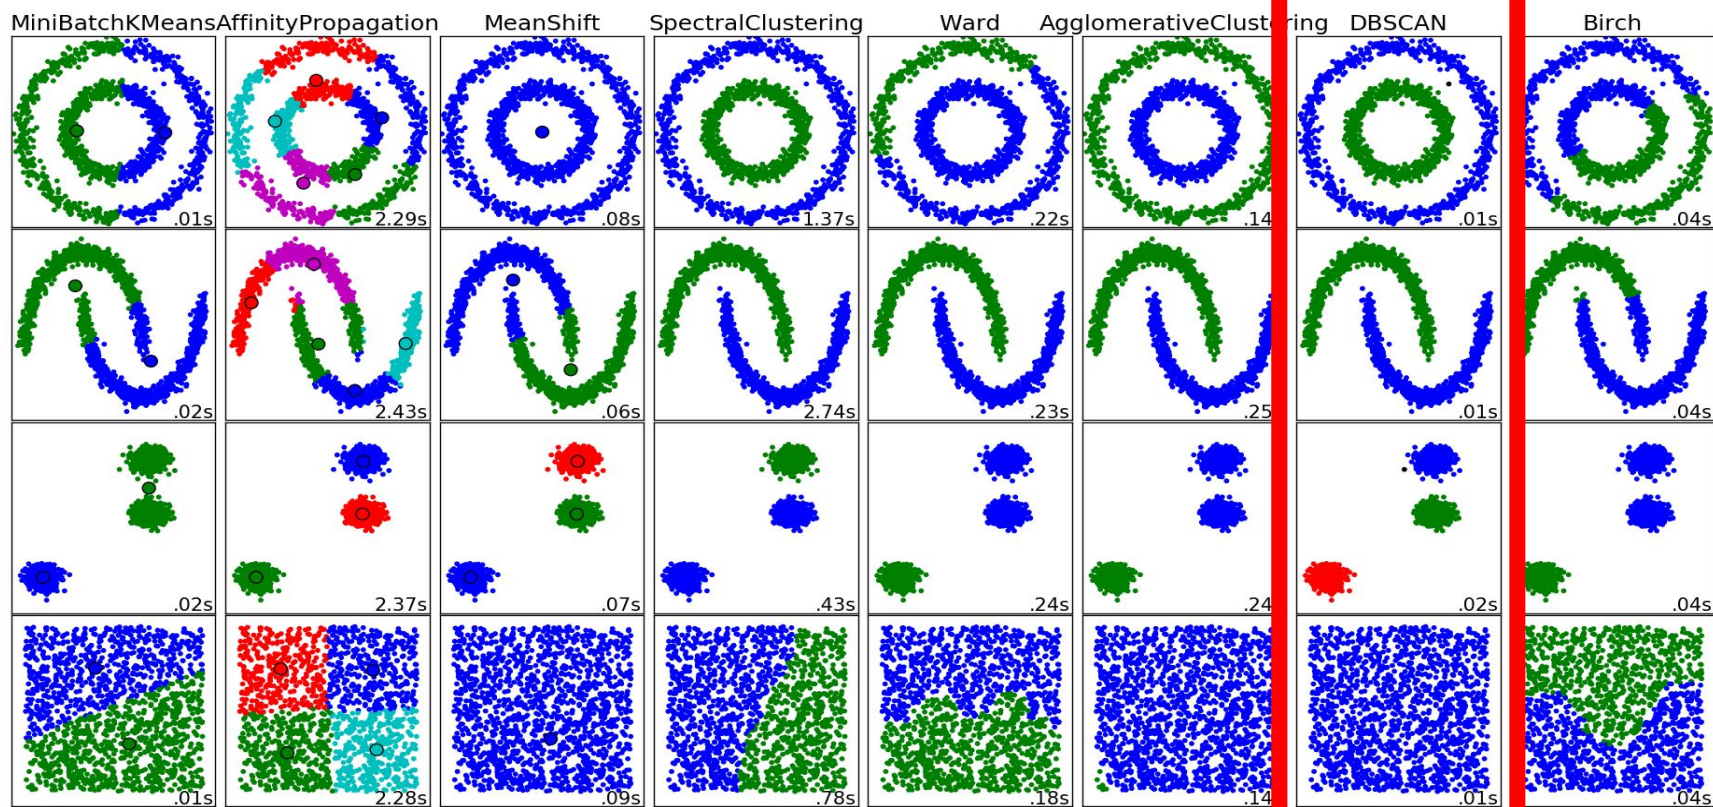
Pros

- Clusters are not necessarily globular
- No choice of number of clusters
- Very efficient implementations exist
- Robust to noise

Cons

- The **eps** and **min_samples** can be hard to tune
- If clusters have significantly different densities it is hard to find a meaningful **eps, min_samples**

Comparisons of clustering algorithms





Hands-on session

01-unsupervised_learning.ipynb